

# Orders of Magnitude:

## Stan Algorithms and Engineering

**Bob Carpenter**

Center for Computational Mathematics

Flatiron Institute

June 2023



# Flatiron Institute

The **mission** of the **Flatiron Institute** is to **advance scientific research** through **computational methods**, including data analysis, theory, modeling and simulation.

- **Center for Computational**: Math (stats/ML), Biology, Neuroscience, Astrophysics, Quantum Physics
- part of **Simons Foundation**
  - US\$300M+ science and education grants per year
  - US\$5B endowment

## Pop Quiz

- Who's the **most famous Stan** in **St. Louis**?

# Stan the Man

(active 1941–1963)

- Not Ulam. Not a stalker fan.
- **Stan “The Man” Musial**. Outfielder, *St. Louis Cardinals*.



*Wikipedia:* One of the **greatest and most consistent** hitters in baseball

# ≡ Inside baseball (metaphor)

---

Article [Talk](#)

---

From Wikipedia, the free encyclopedia

In American slang, the term ***inside baseball*** refers to the minutiae and detailed inner workings of a system that are only interesting to, or appreciated by, experts, insiders, and aficionados.<sup>[1][2]</sup> The phrase was

... ..

# Swing for the fences


- 2011: I move **from industry to Columbia** to work with Gelman.
- I was **getting scooped** on crowdsourcing.
- Michael Collins (CS) suggests I **work on harder problems**.
- I **listened** and we started the **Stan** project.



# Where are the fences?

- Sometimes: an **unknown unknown**.
- **Usually**: about an **order of magnitude** away
- I'm a **computer scientist**, so that's only  $\times 2$

( $\approx$  Gelman & Hill)



**Part I**

**MCMC for Bayes**



# Bayesian Inference

(Bayes  $\approx$  1750, Laplace  $\approx$  1800)

- **unknown** (parameters)  $\theta \in \mathbb{R}^D$ ;    **observed** (data)  $y \in \mathbb{R}^N$

- **Estimation:**

$$\hat{\theta} = \mathbb{E}[\theta \mid y] = \int_{\mathbb{R}^D} \theta \cdot p(\theta \mid y) \, d\theta.$$

- **Event probabilities**

$$\Pr[A \mid y] = \mathbb{E}[\mathbf{I}_A(\theta) \mid y] = \int_{\mathbb{R}^D} \mathbf{I}_A(\theta) \cdot p(\theta \mid y) \, d\theta.$$

- **Posterior prediction**

$$p(\tilde{y} \mid y) = \mathbb{E}[p(\tilde{y} \mid \theta) \mid y] = \int_{\mathbb{R}^D} p(\tilde{y} \mid \theta) \cdot p(\theta \mid y) \, d\theta.$$

# Monte Carlo

(Fermi, von Neumann, Ulam  $\approx$  1930s–1940s)

- Go-to approach to **high-dimensional integration**
- **Randomized algorithm** for deterministic integrals

$$\mathbb{E}[f(\theta) \mid y] \approx \frac{1}{M} \sum_{m=1}^M f(\theta^{(m)})$$

for **independent and identically distributed** (i.i.d.) draws

$$\theta^{(m)} \sim p(\theta \mid y)$$

- Proof: **central limit theorem** (CLT) + law of **unconscious statistician**.

# Markov chain Monte Carlo

(Ulam, late 1940s)

- When **too hard to draw i.i.d.** from posterior.
- Draw  $\theta^{(0)}, \dots, \theta^{(M)}$  from a homogeneous **Markov chain**

$$\theta^{(0)} \sim q_0(\theta) \tag{1}$$

$$\theta_{m+1} \sim q(\theta_{m+1} \mid \theta_m). \tag{2}$$

- **Ergodic theorem** says we can use MCMC just like MC, when chain
  - is **irreducible** (doesn't get stuck),
  - is **aperiodic** (doesn't visit partition cyclically), and
  - has **stationary distribution** s.t.  $\theta^{(m+1)} \sim p(\theta \mid y)$  if  $\theta^{(m)} \sim p(\theta \mid y)$ .
- Proof: equilibrium convergence + law of unconscious statistician

# Accept/reject balance

(Metropolis et al. 1950)

- **propose**  $\theta^* \sim q(\theta \mid \theta_m)$ ,
  - where  $q$  is **symmetric**, i.e.,  $q(\theta' \mid \theta) = q(\theta \mid \theta')$
- **accept** with probability  $\min\left(1, \frac{p(\theta^* \mid y)}{p(\theta_m \mid y)}\right)$

```
theta[0] = q0_rng()
for m in range(M):
    theta_star = q_rng(theta[m])
    u = uniform_rng(0, 1)
    accept = log(u) < log p(theta_star) - log p(theta[m])
    theta[m + 1] = theta_star if accept else theta[m]
```

# Non-reversible proposals

(Hastings 1970)

- **propose**  $\theta^* \sim q(\theta \mid \theta_m)$ ,
  - where  $q$  is **not necessarily symmetric**
- **accept** with probability  $\min\left(1, \frac{p(\theta^* \mid y)}{p(\theta_m \mid y)} \cdot \frac{q(\theta_m \mid \theta^*)}{q(\theta^* \mid \theta_m)}\right) t$
- If  $q$  is symmetric, second term drops out, reduces to Metropolis

## Detailed Balance

- Let  $p(\theta_{m+1} \mid \theta_m)$  be the Markov chain **transition kernel**
  - reject probability makes it mixed discrete/continuous
- Metropolis-Hastings **accept** step ensures MCMC kernel satisfies **detailed balance**,

$$p(\theta \mid y) \cdot p(\theta' \mid \theta) = p(\theta' \mid y) \cdot p(\theta \mid \theta').$$

- ensures chain has stationary distribution  $p(\theta \mid y)$
- given irreducibility and aperiodicity

# HMC

(Duane et al. 1987)

- **couples momentum**  $\rho \in \mathbb{R}^D$  to sample over **phase space**  $\mathbb{R}^D \times \mathbb{R}^D$
- **Hamiltonian**  $H(\theta, \rho) = -\log p(\theta, \rho) = -\log p(\theta) - \log p(\rho)$ ,
  - **Kinetic energy**:  $-\log p(\rho) = -\log \text{normal}(0, I_D) = \frac{1}{2} \cdot \theta^\top \cdot \theta$ .
  - **Potential energy**:  $-\log p(\theta) = -\log p(\theta \mid y)$
- **Hamiltonian Monte Carlo** (HMC) couples two stationary-preserving transition kernels:
  - **Exact momentum** refresh:  $\rho_{m+1} \sim \text{normal}(0, 1)$ .
  - **Metropolis proposal**:  $(\theta^*, -\rho^*)$ , where  $(\theta^*, \rho^*)$  solves Hamiltonian dynamics from initial  $(\theta_m, \rho_{m+1})$  to proposal  $(\theta^*, \rho^*)$  at time  $t$ .
  - Metropolis **corrects numerical integration error** solving ODE
  - Momentum **flip for reversibility** (required, but erased)


# Why is HMC so good?

- **Long distance** proposals with **high acceptance** rates
- **Effective sample size** (and hence **mixing**) is proportional to expected squared jump distance
- An **exact solution** preserves initial Hamiltonian, so 100% accept
- The **leapfrog integrator** used to solve dynamics is **symplectic**
  - steps preserve volume (hence no Hastings correction for Jacobian)
  - symplectic integrators very good at **preserving Hamiltonian**
  - not so great at **solving dynamics**, but doesn't matter—we only need **long jumps with high acceptance**



# HMC is hard to tune

- **y-axis:** ESS;      **x-axis:** step size;      **facets:** steps
- **blue:**  $Y$ ;      **red:**  $Y^2$
- **dashed lines:** NUTS



## HMC + Euclidean metric

- add symmetric, positive-definite **metric**  $M$  to define distance

$$||\theta|| = \theta^\top \cdot M \cdot \theta \quad (3)$$

$$d(\theta, \theta') = ||\theta - \theta'|| \quad (4)$$

- Stan estimates  $M$  as inverse posterior covariance,

$$\widehat{M} \approx \text{cov}[\theta \mid y].$$


- **kinetic energy** now  $-\log p(\rho) = \frac{1}{2} \cdot \rho^\top \cdot M \cdot \rho,$
- so **momentum refresh**  $\rho \sim \text{normal}(0, M^{-1}).$

## No-U-turn sampler (NUTS) (Hoffman and Gelman 2013)

- **Tuning HMC** dynamics (step size, number of steps) is **very hard**
- No-U-turn sampler (NUTS) automatically tunes
  - **stepsize** during warmup iterations
  - **metric** during warmup iterations
  - **number of steps** dynamically
- Betancourt (2017) added several improvements to NUTS for Stan

## NUTS (cont.)

- Sampling algorithm
  - randomly simulate **forward vs. backward** time, **doubling steps**
  - **until U-turn** (momentum brings ends closer)
  - **select a step** along path, **biased** toward **last doubling**
  - **slice sampling**, revised to more efficient **multinomial**



# Generalized HMC

(Horowitz 1991)

- Uses **partial momentum refresh**
  - preserves momentum for directed movement across posterior
- Fix  $\lambda \in (0, 1)$  (lower preserves more momentum)

- **G-HMC refresh:**

$$\rho_{m+1} = \sqrt{\lambda} \cdot z_m + \sqrt{1 - \lambda} \cdot \rho_m,$$

where  $z_m \sim \text{normal}(0, M^{-1})$ .

- This is also an **exact update**
  - i.e., if  $\rho_m \sim \text{normal}(0, M^{-1})$ , then  $\rho_{m+1} \sim \text{normal}(0, M^{-1})$ .


## Uh oh! What about the Flip?

- **Why not** use G-HMC with a single step and tune  $\lambda$ ?
- Remember that **momentum flip**?
  - it's required for **reversibility** of Metropolis
  - **thrown away** in basic HMC by composing momentum update
  - but **preserved** in G-HMC
- **Without high acceptance**, momentum flip produces **random walk**.
- High acceptance means **small step size** in Hamiltonian dynamics.

# Non-reversible accept

(Neal 2020)

- Instead of uniform  $u$ , use **sawtooth** pattern, **jittered** for ergodicity
  - **not reversible**, but **preserves stationary**
- Iteration vs. accept probability  $u$  (red reject) **groups acceptance**:



# Delayed Rejection

(Mira 2001)


- $K$ -step delayed rejection involves  $K$  distinct proposals
  - Step 1. Propose and accept/reject as usual with Metropolis.
  - Step 2. **If rejected, propose again** with a new proposal and accept/reject with **Metropolis-Hastings**
  - $\vdots$
  - Step  $K$ . If rejected, propose one last time.
- Need **Hastings correction** to ensure detailed balance.
- Proposals may depend on previous proposal(s).
- **Innovation:** Only retry if acceptance probability was low



# Delayed Rejection HMC

(Modi et al. 2022)

- **Multiple scales** (varying curvature) along an **entire path**
- Each **retry cuts step size** by constant  $c$  and **multiplies steps** by  $c$  (e.g.,  $c = 2$  or  $c = 5$ )
  - earlier attempts kept step size and extended path
  - ours better if gradient-based **Hamiltonian diverged**



(funnel draws/accept)

# MEADS

(Hoffman and Sountsov 2022)

- **Massively parallel** version of Neal's non-reversible accept **G-HMC**
- Uses **complementary chains for adaptation**
  - novel, efficient **principal eigenvalue** estimation of step size
  - covariance in complementary chains used to **estimate metric**
  - **accelerates adaptation**, robust to single chains getting stuck (Bales 2019)
  - much easier to **parallelize**
  - **little waste** and **easy restart**
- conveniently a Markov chain **w.o. warmup phase**
  - still need **burn-in**

# DR-G-HMC

(Modi, Roualdes, ... in progress)

- Apply **delayed rejection** to **generalized HMC**
- **Retries** use **one step** instead of constant time
- **Multiple scales** adapted **within a trajectory**
- solves **non-centered funnel** (like DR-HMC)
- Use MEADS-like **parallel tuning**

# Convergence: nested $\hat{R}$

(Margossian et al. 2022)

- **GPUs** run **1000+** parallel chains
- **One draw/chain minizes wall time**
- **Nest blocks** with **same init**; monitor **transient bias + variance**

$$\begin{array}{c} \text{Squared Error} = \underbrace{\text{Squared Bias} + \text{Non-stationary Variance}}_{\text{Decays as the chains converge}} + \underbrace{\text{Persistent Variance}}_{\substack{\text{Decays as the numk} \\ \text{subchains increases}}} \end{array}$$

Monitored by  $n\hat{R}$

# **Part II**

## **Stan Language**

# Tuples

(next release)

- Required **design doc** choices
  - mainly on syntax and including **named structs**
- **Massive refactoring** required in transpiler & I/O (thanks Brian and Steve)
  - **my bad** in original code assuming **dense rectangular**

```
tuple(int, real) x = (42, 3.14); // type & construct
```

```
int a = x.1;    real b = x.2;    // rvalue (access)
```

```
x.1 += 1;      x.2 = sqrt(x.2); // lvalue (assign)
```

# Closures and lambdas

(design accepted)

- **Closures bind** non-local variables in function bodies
- **Lambdas** for anonymous inline function definitions (C++-like syntax)
- Full **type** support in the object language
  - variables with function types and assignment
  - function arguments with function types (i.e., higher-order functions)
- Support **comprehensions** with **partial evaluation**
  - e.g., scalable GP kernels

# Ragged arrays

(design pending)

- Not much more to say about this
- Challenging to **statically size**
- Need to **generalize math** library
- **Bigger challenge: sparse matrices**
  - Eigen C++ support is poor
  - Have to roll our own algorithms (e.g., log determinant)



## More types and constraints

- **Orthonormal matrices** (generalizes unit vector)
  - e.g., hyperspherical statistics and rotations
  - tricky  $SO(N)$  geometry (rotation vs. reflection)
- **Composable** transforms
  - e.g., affine compositions for hierarchical rates or probabilities
- **Pluggable** transforms
  - e.g., specialized simplex transforms for scale

# **Part III**

## **Automatic Differentiation**

# Mat-var to Var-mat

(underway)

- **Reverse-mode** autodiff uses **autodiff variables**
  - one per value evaluated in the constraints and log density
  - C++ class instance store a **value** and **adjoint**
  - virtual function to apply chain rule
- Stan matrices currently store a **matrix of variables**
- Moving to more efficient **variable of matrixes**
  - more **memory locality**
  - less **copying**
- **huge refactor** of math lib and transpiler

# Expression templates

(well underway)

- original Stan **evaluated** expression templates
- goal: **lazy evaluation** plus **static compilation**
- improved Stan tries to **propagate** expression templates
  - **reduces copying**
  - accelerate **composed code compilation** avoids eager, unnecessary copies
  - Fortrain-like speed; how C++ is faster than C
- very **tricky C++**

# GPU kernel

(starting)

- Very **expensive to move** data between GPU and CPU
- now use where order of computation */gg* order of data
  - e.g., matrix-matrix multiply, Cholesky factorization, **not** matrix-vector multiply, dot products, etc.
- Write **GPU code** for entire **math library**
- **Order of magnitude (or two!)** speedup (cf. JAX)
- **Enormous** undertaking

# Part IV

## Approximate Inference

(cf. **MCMC**, which is asymptotically **exact**)

# Laplace Approximation

(released)

- Given a **posterior mode**

$$\theta^* = \arg \max_{\theta} p(\theta | y)$$

- Second-order Taylor expansion** is called **Laplace approximation**

$$\theta \sim \text{multi\_normal}(\theta^*, (-H)^{-1}(\theta^*)),$$

where  $(-H)^{-1}$  is the inverse negative Hessian

- pairs with our built-in gradient-based **quasi-Newton optimization**
  - uses **L-BFGS** for local curvature adjustment (i.e., **precondition**)

# Autodiff Variational Inference (ADVI)

- Given posterior  $p(\theta | y)$  on the **unconstrained scale**
  - i.e.,  $p(\theta | y) > 0$  for all  $\theta \in \mathbb{R}^D$
- **Variational approximation** is  $\text{multi\_normal}(\theta | \hat{\mu}, \hat{\Sigma})$ , where

$$\hat{\mu}, \hat{\Sigma} = \operatorname{argmin}_{\mu, \Sigma} \text{KL}[\text{multi\_normal}(\theta | \mu, \Sigma) || p(\theta | y)]$$

- **Sample**  $\theta^{(1)}, \dots, \theta^{(M)} \sim \text{multi\_normal}(\theta | \hat{\mu}, \hat{\Sigma})$ ,
- **Inverse transform** draws back to constrained scale
- **Improvements** by (1) Welandawe, Andersen, Vehtari, and Huggins (2022); (2) Domke and Agrawal (2022).



## ADVI objective evaluation

- KL-divergence is integral, w.  $q(\theta) = \text{multi\_normal}(\mu, \Sigma)$ .

$$\text{KL}[q \parallel p] = \underbrace{\int_{\mathbb{R}^D} q(\theta) \cdot \log q(\theta) \, d\theta}_{\text{entropy of } q} - \underbrace{\int_{\mathbb{R}^D} q(\theta) \cdot \log p(\theta \mid y) \, d\theta}_{\text{cross entropy } q \text{ to } p} \quad (5)$$

$$\approx -H[\text{multi\_normal}(\mu, \Sigma)] + \frac{1}{M} \sum_{m=1}^M \log p(\theta^{(m)} \mid y), \quad (6)$$

where  $\theta^{(m)} \sim \text{multi\_normal}(\mu, \Sigma)$ .

# ADVI SGD

- Need **gradient**  $\nabla_{\theta} \text{KL}[q || p]$  to minimize

$$\arg \min_{\mu, \Sigma} \text{KL}[\text{multi\_normal}(\theta | \mu, \Sigma) || p(\theta | y)]$$

- **Entropy** term can be handled analytically or by Monte Carlo
- **Cross-entropy** requires **stochastic gradient**,

$$\nabla_{\theta} \sum_{m=1}^M \log p(\theta^{(m)} | y) = \sum_{m=1}^M \nabla_{\theta} \log p(\theta^{(m)} | y),$$

with nested derivative by **automatic differentiation**

# ADVI improvements

(in progress)

- Need to select a **step size**
  - current algorithm too weak; parallel grid search way better
- Need to select an **SGD algorithm**
  - ADVI is vanilla SGD; **ADAM**'s persistent momentum better
- Need to select **gradient estimator**
  - Stan uses vanilla reparameterization; **stick the landing** is better
- Stan uses standard normal **initialization**
  - **Multiple inits** better (Laplace, prior, standard normal, etc.)
- Stan just **returns constrained draws**
  - **importance resampling** is better







# Pathfinder

(next release!)

- **Quasi-Newton** variational inference with **L-BFGS**
- At each point on optimization trajectory:
  - lay down a multivariate normal approximation
  - covariance is approximate **negative inverse Hessian**
  - estimated with finite differences of autodiff gradient
  - low-rank (number of gradients  $J \approx 5$ ) plus diagonal:  $\mathcal{O}(D \cdot J^2)$
- Choose approximation with **lowest KL-divergence** (Monte Carlo)
- Importance **resample**
- More robust: **multiple paths**, combined importance resampling

# Pathfinder illustrated

- Ellipses = Taylor approximation; **lower left** is best (but **overconcentrated**)



# The future: normalizing flows (Domke and Agrawal)

- Like ADVI with **normalizing flows** as approximating family
  - Generate **inital**  $X \sim q_X(x)$  (cf. Domke and Agrawal)
  - **Transform** with **deep neural net**  $f_\beta : \mathbb{R}^D \rightarrow \mathbb{R}^D$ ,  $\Theta = f_\beta(X)$ ,

$$q_\Theta(\theta \mid \beta) = q_X(f_\beta^{-1}(\theta)) \cdot \left| J_{f_\beta^{-1}}(\theta) \right|$$

- **Optimize** as in ADVI:  $\beta^* = \arg \min_\beta \text{KL}[q_\Theta(\theta \mid \beta) \parallel p(\theta \mid y)]$
- For VI, neural net must have **efficient Jacobian and sampling**
- Real non-volume preserving (**RealNVP**) flows work with **JAX**
  - complementary affine layers,  $\tanh()$  non-linearity ( $\approx 10$  deep, 12 wide)
  - fits centered hieararchical IRT-2PL (additive + multiplicative + funnel)

## More work in progress

- score-function based **control variates**
- **amortized variational inference** (BRMS)
- **black-box nested Laplace** approximation
- low-dimensional (1–2), massively **scalable Gaussian processes**